

NTURACING

MCU Programming

W2: Program Execution

by 羅紀翔

2024-25 FSAE Season



Program
Execution

1

2

C Technique 2
Macros

Exceptions

3

4

LAB 2
Debug

Program Execution



References: CH4, 7, 8, *The Definitive Guide*; Ch 20, 22, *Mastering STM32*; Stack Memory, Code Inside Out:

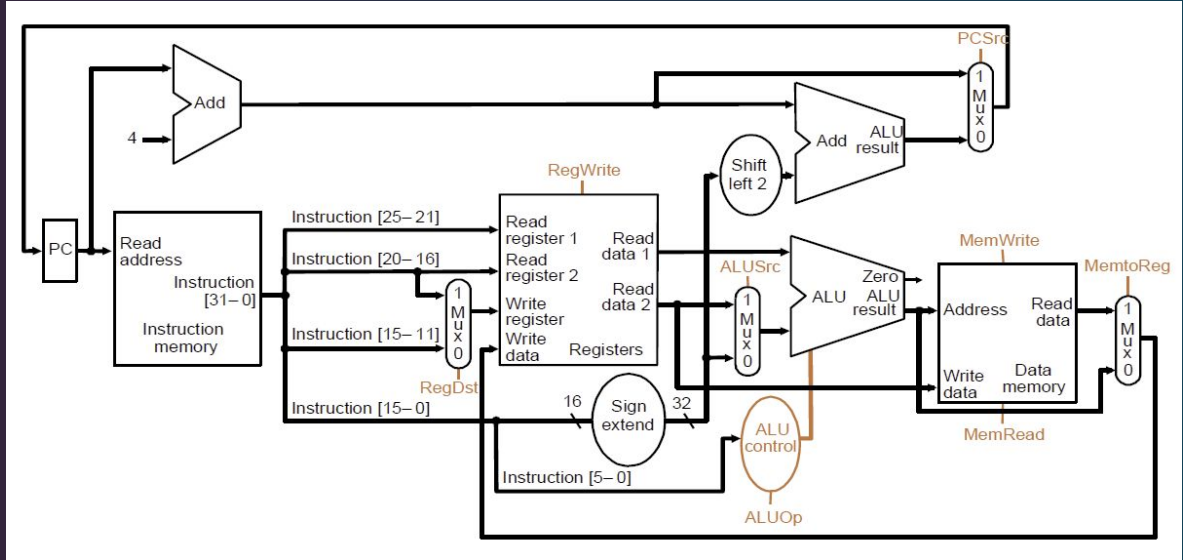
<https://codeinsideout.com/blog/stm32/stack-memory/>

CPU Instruction

- Essential instructions:
 - **I/O**: Move data from/to memory
 - **Arithmetic**: Addition, subtraction, ...
 - **Logical**: AND, OR, NOT, ...
 - **Control flow**: Jump, branch, ...
- Instruction can be written by assembly
- Basic syntax: [label] mnemonic [operands]
- Checkout Compiler Explore (<https://godbolt.org/>) to see what assembly does the source code compiles to

```
.LC0:
    .ascii  "hello world\000"
main:
    push    {fp, lr}
    add     fp, sp, #4
    ldr     r0, .L3
    bl      puts
    mov     r3, #0
    mov     r0, r3
    sub     sp, fp, #4
    pop     {fp, lr}
    bx      lr
.L3:
    .word   .LC0
```

- Single cycle MIPS CPU
- Fetch, decode, execute cycle
- Take **Computer Architecture** to learn more




Features for High-Level Languages

- "A programming language with strong abstraction from the details of the computer."
(from Wikipedia: High-level programming language)
 - C is just a thin layer of abstraction over assembly
- Automatic memory management: stack, heap (only stack in C)
- Function call
- Runtime
 - Interpretation, VM
 - Exception
 - Library
 - C requires data and zero initialization

Execution

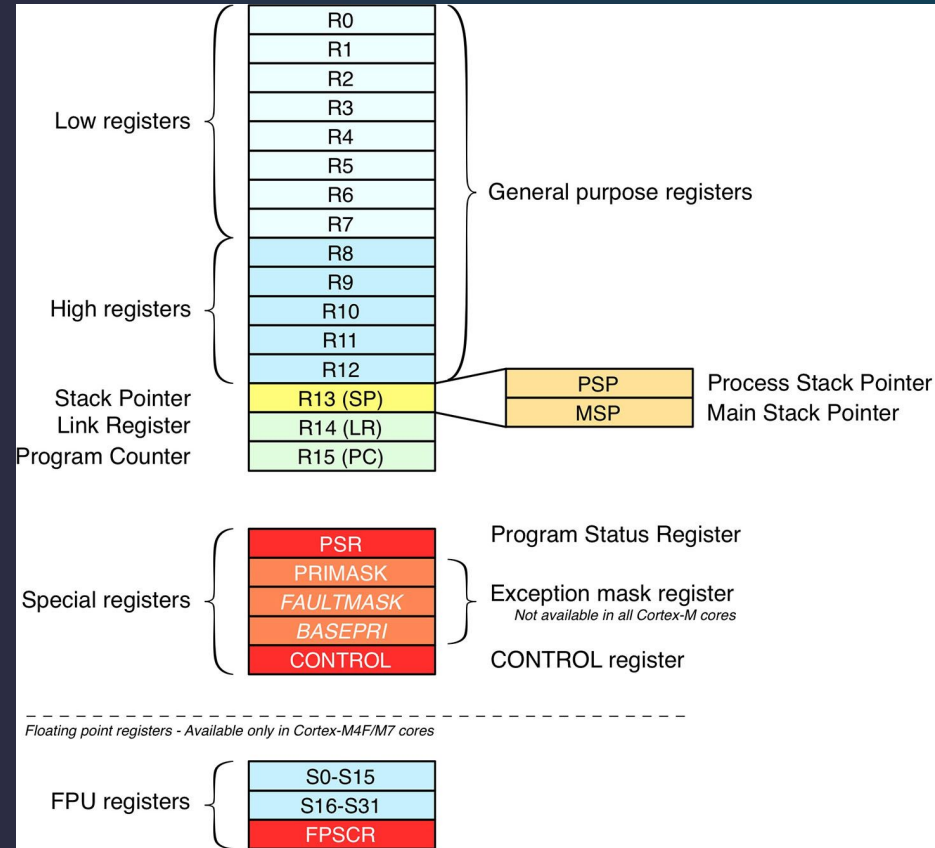
Stack

- Stores intermediate values in execution
- Stores function arguments
- Saves execution context when procedure (function) call or during ISR
- Cortex-M is full descending stack



Recall: Registers

- Frame pointer (**FP**): R7, if used (typically), stores the start of the function frame
- Stack pointer (**SP**): the top of the stack
 - **Shadowed** (Banked), usually used by embedded OS to differentiate privileged and unprivileged stacks
- Link register (**LR**): return address of functions
- Program counter (**PC**): location of the next instruction



Function Call

- Calling convention, aka Application binary interface (**ABI**) (c.f. API)
- ARM Architecture Procedure Call Standard (AAPCS)
- **Caller-saved** : R0~R3, R12, can be freely used by functions
 - R12 is caller-saved since it's used as scratch register that the compiler can use it to store temporary values and does not need to save it to the stack
- **Callee-saved** : other registers, must be saved (typically in the stack) before used
- Arguments: R0~R3
- Return value: R0 (and R1 if return value is 64-bit)
- Ref. Calling convention, Wiki: [https://en.wikipedia.org/wiki/Calling_convention#ARM_\(A32\)](https://en.wikipedia.org/wiki/Calling_convention#ARM_(A32))

Function Prologue and Epilogue

- Procedures to conform to calling convention
- Prologue
 - Push LR and FP
 - Push callee-saved registers
 - Allocate stack for local variables
- Epilogue
 - Deallocate stack
 - Pop callee-saved registers
 - Pop FP and LR
- Overhead of function calls!
 - Also heavily optimized in modern CPU and compiler

Execution

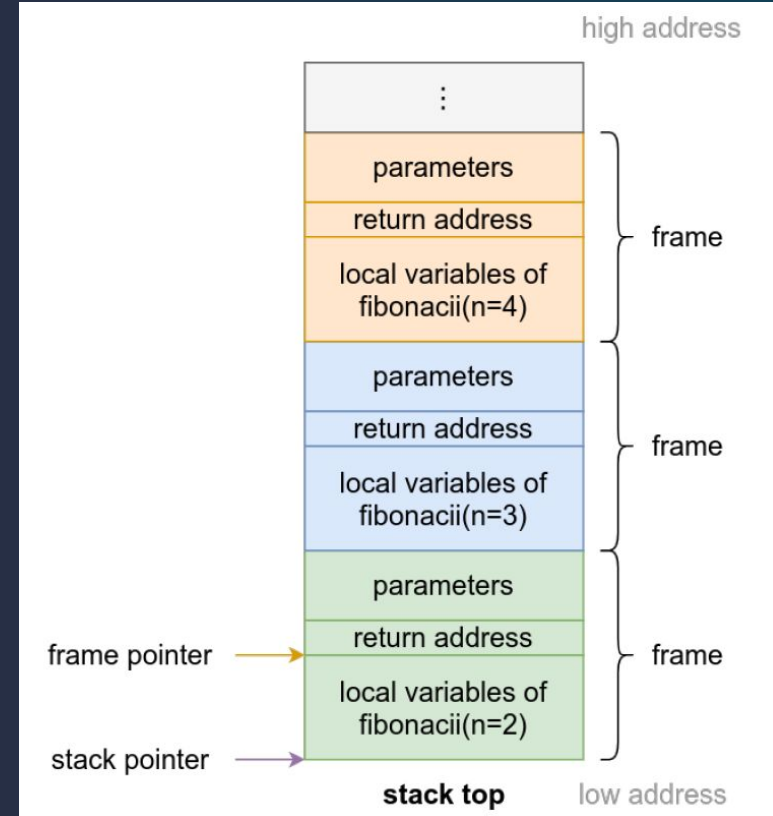
```
int add(int a1, int a2, int a3, int a4, int a5) {  
    return a1 + a2 + a3 + a4 + a5;  
}  
  
int five() {  
    return add(1, 1, 1, 1, 1);  
}
```

```
five:  
  
    push    {fp, lr}  
    add     fp, sp, #4  
    sub     sp, sp, #8  
  
    mov     r3, #1  
    str     r3, [sp]          ; passing a5 via stack  
    mov     r3, #1  
    mov     r2, #1  
    mov     r1, #1  
    mov     r0, #1  
    bl      add  
    mov     r3, r0  
  
    mov     r0, r3  
    sub     sp, fp, #4  
    pop     {fp, lr}  
    bx      lr
```

```
add:  
  
    push    {fp}              ; push fp  
    add     fp, sp, #0        ; set fp to the location of  
                                ; previous fp  
    sub     sp, sp, #20       ; allocate 20 bytes of stack  
    str     r0, [fp, #-8]     ; save used registers to  
  
stack  
  
    str     r1, [fp, #-12]  
    str     r2, [fp, #-16]  
    str     r3, [fp, #-20]  
  
    ldr     r2, [fp, #-8]     ; load a1-a4 from stack  
    ldr     r3, [fp, #-12]  
    add     r2, r2, r3  
    ldr     r3, [fp, #-16]  
    add     r2, r2, r3  
    ldr     r3, [fp, #-20]  
    add     r2, r2, r3  
    ldr     r3, [fp, #4]      ; load a5 from stack  
    add     r3, r2, r3  
  
    mov     r0, r3  
    add     sp, fp, #0        ; restore original sp from fp  
    pop     {fp}  
    bx      lr                ; return
```

Function Frame

- Stack frame saves the currently running function
- Infinite recursion will cause stack overflow
 - You can also see the overhead of recursion
- FP in the stack can be used to debug function call stack
- FP is optional as compiler knows how many stack the function requires, could be optimized via **frame pointer omission**, however, function call stack can't be debugged in this case



Bootup Sequence - Hardware

- **Reset**: Set all CPU and peripherals registers to predefined values
- After reset: Cortex-M loads SP and PC from 0x00000000 and start execution
- STM32 uses hardware remap to map 0x00000000 to different addresses depending on boot select pins:

Boot mode selection pins		Boot mode	Aliasing
BOOT1	BOOT0		
x	0	Main Flash memory	Main Flash memory is selected as the boot space
0	1	System memory	System memory is selected as the boot space
1	1	Embedded SRAM	Embedded SRAM is selected as the boot space

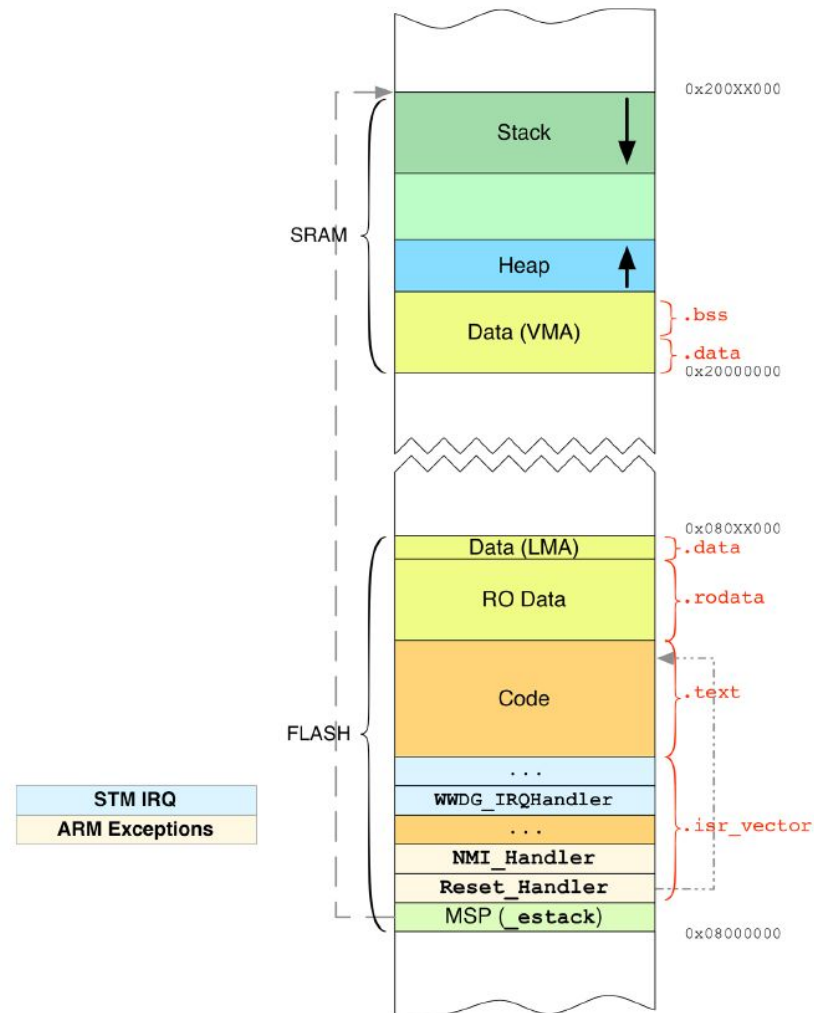
- Flash: 0x08000000
- ST's bootloader: 0x1fff0000
- SRAM: 0x20000000

Execution

Recall: Program in Memory Layout

- Memory sections
 - **.text**: instructions
 - **.data**: initialized data
 - **.bss** (blocks started by symbols): uninitialized data
 - **.rodata**: const data and strings
- **.isr_vector**:
 - Entry point for exceptions
- Startup code
 - **Data initialization**: .data is copied from flash to SRAM
 - **Zero initialization**: .bss is set to all zeros

Question: How startup code works?



Bootup Sequence - Setup C Runtime

- Referred as crt0 in Linux (<https://en.wikipedia.org/wiki/Crt0>)
- Different frameworks/ OSes have different implementations, here we discuss STM32CubeIDE
- Recall: Linker script
 - Refer to *_FLASH.ld in STM32CubeIDE
 - Defines _sdata, _edata, etc. to mark the ranges of some sections
- Recall: PC is loaded from the second entry of isr_vector, which is **Reset_Handler** and starts execution from there
- Startup code (Reset_Handler)
 - Written in assembly as C environment is not ready
 - Performs data and zero initialization and jump to main()
 - Refer to start_*.s in Core/Startup/ in STM32CubeIDE

C Technique 2: Macros



References: Macro Metaprogramming
(<https://mailund.dk/posts/macro-metaprogramming/>)

Macros

- #define
- # (stringizing operator)
- ## (token pasting operator)
- Use examples from:
 - Basics: Macro Metaprogramming (<https://mailund.dk/posts/macro-metaprogramming/>)
 - Advanced: C Preprocessor tricks, tips, and idioms (<https://github.com/pfultz2/Cloak/wiki/C-Preprocessor-tricks,-tips,-and-idioms>)
 - Zephyr RTOS utilities macros (https://docs.zephyrproject.org/latest/doxxygen/html/group__sys-util.html)

Exceptions



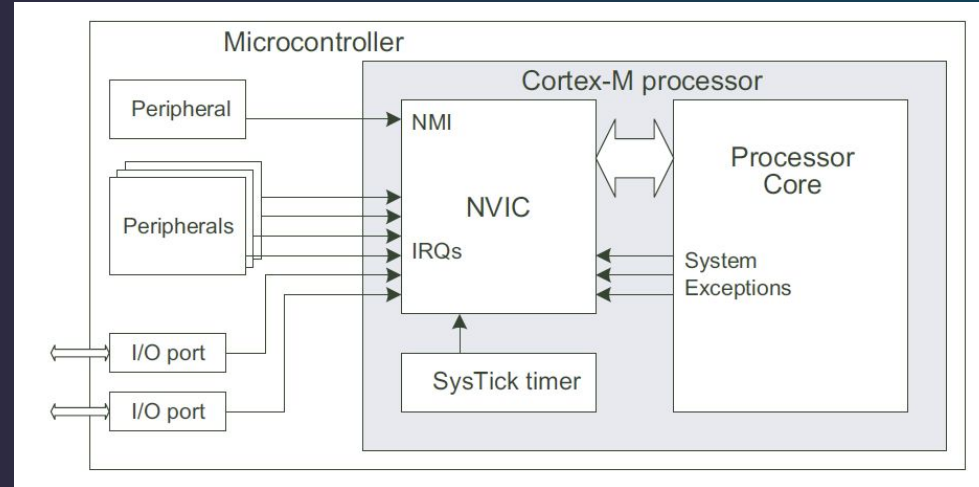
References: CH7, 8, *The Definitive Guide*; Ch7, *Mastering STM32*

Recall: Exceptions

- Hardware: interrupt, software: fault
- Suspends current CPU execution and enters interrupt service routine (ISR), functions to handle exceptions
- Various types to represent different events

Nested Vectored Interrupt Controller (NVIC)

- Interrupt controller in Cortex-M
- **Nested**: An exception can be handled while another is executing
- **Vectored**: Supports multiple exceptions by a vector table
- Exception number 1~15 are system exception and 16 and above for external interrupt used by peripherals via Interrupt Request (**IRQ**) lines



isr_vector

- Entry points for each exceptions
- Typically starts at 0x00000000, can be relocated by modifying VTOR register of NVIC
- In practice, it's an array of function pointers of prototype void isr_func(void)
- Defined in start_*.s in Core/Startup/ in STM32CubeIDE
- The first section in flash, defined in *_FLASH.ld in STM32CubeIDE

Number	Exception type	Priority ^a	Function
1	Reset	-3	Reset
2	NMI	-2	Non-Maskable Interrupt
3	Hard Fault	-1	All classes of Fault, when the fault cannot activate because of priority or the Configurable Fault handler has been disabled.
4	Memory Management ^c	Configurable ^b	MPU mismatch, including access violation and no match. This is used even if the MPU is disabled or not present.
5	Bus Fault ^c	Configurable	Pre-fetch fault, memory access fault, and other address/memory related.
6	Usage Fault ^c	Configurable	Usage fault, such as Undefined instruction executed or illegal state transition attempt.
7	SecureFault ^d	Configurable	SecureFault is available when the CPU runs in <i>Secure state</i> . It is triggered by the various security checks that are performed. For example, when jumping from Non-secure code to an address in Secure code that is not marked as a valid entry point.
8-10	-	-	RESERVED
11	SVCall	Configurable	System service call with SVC instruction.
12	Debug Monitor ^c	Configurable	Debug monitor – for software based debug.
13	-	-	RESERVED
14	PendSV	Configurable	Pending request for system service.
15	SysTick	Configurable	System tick timer has fired.
16-	IRQ	Configurable	IRQ Input

[47/239/479]^e

Exception States


- Exceptions have three binary states
 - Disabled (default) or enabled
 - Pending (waiting to be handled) or not pending
 - Active (being handled) or inactive
- Exceptions will only be handled if it's enabled and pending, then it will enter active state

Exception Masking

- Exceptions can be temporarily disabled (called **masking**): `__disable_irq()`, `__enable_irq()`
- Typically used by embedded OSes to perform critical sections
- Pending state is ignored if it's masked (but it will be handled right after when it's unmasked)
- Reset, NMI, hard fault, and SVCcall are not maskable

Exception Priority

- A higher-priority exception (smaller number in priority level) **preempts** a lower-priority exception.
- Priority levels are up to 8 bits (4 bits in STM32) and split into **preemption priority** and **sub priority** of configurable lengths.
 - Preemption Priority: Higher-priority ones preempt lower-priority ones.
 - Sub Priority: Higher-priority ones will run first when fired at the same time.



Exception Execution

- **Stacking** : When exceptions are handled, the caller-saved registers (r0~r3, r13) and LR and other control registers are saved to the stack, they will be **unstack** ed after ISR had finished so that the registers are the same before and after an exception is handled
- While some CPUs uses special instructions to return from ISR (which requires special support from the compiler or inline assembly), Cortex-M uses a special **EXC_RETURN** value loaded to LR to indicate returning form ISR
- Some optimizations:
 - **Tail-Chaining** : No stacking/unstacking when exceptions are handled in series
 - **Late Arrival** : Exceptions with higher priority is fired during stacking, then it will be handled first without additional stacking

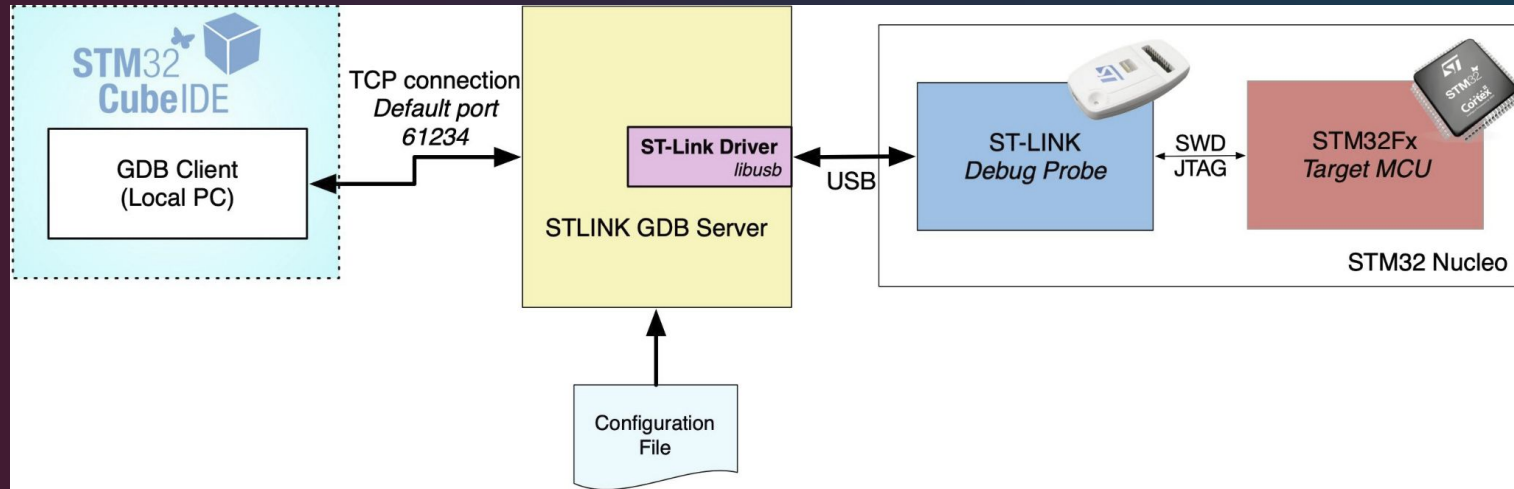
LAB 2: Debug

References: Ch 5, *Mastering STM32*



Debug Pipeline

- Debug communication protocols:
 - Joint test action group (**JTAG**)
 - Serial wire debug (**SWD**) (only for Cortex-M)



Debug Features

- Step over, step into, step out
- Breakpoints
- Variables
- Function call stack
- Memory
- Peripheral registers
- **Cotrex-M live watch**